

# CS6400-A Project Proposal (Group 4): Hybrid Retrieval with Budgeted Evidence Selection

Navneet  
Georgia Institute of Technology  
n41@gatech.edu

Jigyasa Hasija  
Georgia Institute of Technology  
jhasija3@gatech.edu

Joseph Souaiby  
Georgia Institute of Technology  
jsouaiby3@gatech.edu

## ABSTRACT

We will implement a RAG system with hybrid BM25 and dense retrieval, pairwise cross-encoder reranking, and section-aware evidence selection under a token budget. On identical chunks, we will compare it with both required baselines on QASPER [7] and MuSiQue [25]. Our original implementation and controlled ablations will measure packed gold-evidence recall, reranking pool-size trade-offs, section affinity, and system cost. Retrieval will make no generative LLM calls; a fixed generator will produce answers with provenance for supplementary QA evaluation. We will report negative results as well as gains.

## 1 INTRODUCTION

A Retrieval-Augmented Generation (RAG) system [14] answers a question in two stages: retrieval finds the passages most likely to contain the answer, and generation reads those passages and writes a response. On long documents such as research papers, the generator can only use what fits in its context, so retrieval quality should be judged on the packed context under a token budget, not only on the ranked list [2].

**What is known.** BM25 [21] matches exact keywords, while dense retrieval [10] matches meaning. Combining the two, known as hybrid retrieval, can outperform either method alone. The two rankings are usually merged either by Reciprocal Rank Fusion (RRF) [6], which combines ranks, or by a convex combination, which takes a weighted sum of normalized scores. When its weight is tuned, the convex combination matches or beats RRF [4]. Cross-encoder reranking [16] was the largest single gain on all three benchmarks of Singh et al. [24], and BM25 interpolation can improve dense-retriever scores [26]; this does not establish a benefit after cross-encoder reranking. Visconde reranks every paragraph of a QASPER paper [18]. Most directly, Singh et al. [24] held a strong listwise reranker fixed on QASPER and MuSiQue and found that rank fusion, RAPTOR-style hierarchy [22], and graph expansion then add no reliable gain. Structure-aware methods retrieve larger units or neighbouring chunks [9, 13, 15]. Budgeted packing has been studied with coverage objectives on multi-hop data [2] and with context-budget diagnostics on QASPER [11].

**The gap.** Our study will examine four connected questions for QASPER:

- (1) How much gold evidence reaches the generator at a fixed token budget after each stage. Singh et al. [24] report ranking and answer metrics, including nDCG@10 and Success@10, but do not isolate paragraph-evidence survival through our budgeted packing stages.
- (2) Whether their conclusion holds for a pairwise cross-encoder and for weaker BM25 and MiniLM [23] first stages, both of which they explicitly leave to future work [24].

- (3) How small the reranking pool can be before evidence is lost.
- (4) Whether section-aware selection still helps once a reranker has ordered the paragraphs [13, 15].

QASPER’s section structure [7] motivates testing section affinity. We will measure the proportion of multi-evidence questions and the degree of evidence co-location on dev before evaluating this hypothesis.

**What we will do.** We will build the pipeline with every stage switchable, using the required baselines on identical chunks. We will score each configuration by gold-evidence recall and all-evidence coverage of the packed context across budgets of 256–2048 tokens. We will record candidate-pool misses and selection losses separately. For selection losses, logs will retain score order, token costs, and the budget state to diagnose ordering and fit failures without asserting a causal partition. We will address three research questions:

- **RQ1 (first stage and pool).** How much gold evidence does each first stage (BM25, dense, RRF [6], convex [4]) deliver to the packed context after pairwise reranking [16], and how small can the pool  $N$  be before packed recall drops? Once every paragraph is scored, first-stage pool recall is no longer limiting; the pool-size trade-off will be our focus.
- **RQ2 (selection).** At the same budget, does section-aware selection raise all-evidence coverage on multi-evidence questions over top- $k$  after reranking, and at what cost to single-evidence questions? We will compare it against redundancy-aware selection [1, 5, 17, 27]. On multi-hop Wikipedia data, coverage-based packing helps [2], but on QASPER a diversity penalty could exclude complementary evidence from one section. We will test both possibilities.
- **RQ3 (features).** Do BM25 and dense scores still add to cross-encoder relevance as selector features? Prior work motivates fusion before reranking [26]; whether these features help after reranking is an empirical question.


A negative answer to any question would be a reportable result.

## 2 PROPOSED APPROACH

The pipeline (Figure 1) will retrieve broadly with a hybrid of BM25 and dense retrieval, score candidates precisely with a cross-encoder, and then pack an evidence set into a fixed token budget. Every stage will be switchable; the required baselines will be the first-stage boxes alone, packed with top- $k$ .

### 2.1 End-to-End Design

**Parsing.** QASPER [7] ships each paper as a title, an abstract, and full text organized into named sections of paragraphs, plus figures,



**Figure 1: System overview.** Shaded blue boxes are the stages we will contribute; gray boxes are the required baselines, which will be evaluated along with top- $k$  packing under the same budget. Gold-evidence recall will be measured on the selector output, i.e., on what the generator actually receives.

tables, questions, answers, and evidence annotations. We will read this structure directly; no PDF parsing will be needed.

**Chunking.** One chunk will be one paragraph, the same granularity as QASPER’s gold evidence. Gold strings will be matched exactly after Unicode and whitespace normalization. Unmatched evidence will be retained as unreachable gold IDs in every method’s denominator; no fuzzy matching will be used. We will canonicalize repeated normalized paragraphs and report ambiguous and unmatched evidence separately; no method will use test evidence in retrieval. Each chunk will keep a stable ID, its section title, and its token count.

**Indexing.** We will implement a BM25 [21] inverted index ourselves, and build a dense index of all-MiniLM-L6-v2 [20, 23] embeddings with exact cosine search. MiniLM reads at most 256 wordpieces in its sentence-transformers configuration [23], and may truncate long paragraphs. We will report truncation rates for all paragraphs and gold paragraphs; every dense and hybrid configuration will use the same input limit and retain full source paragraphs for packing. For per-paper retrieval, BM25 IDF will be computed over the training corpus rather than one paper, because term rarity inside a single paper of a few dozen paragraphs is a weak signal.

**Hybrid first stage.** BM25 and dense rankings will be fused by Reciprocal Rank Fusion [6] (RRF) and, separately, by a convex combination of min-max normalized scores [4]. The top  $N$  candidates will form the reranking pool;  $N$  will be an experimental variable. RRF will use  $1/(60 + \text{rank})$  per channel. Constant-score channels will normalize to zero; ties will be broken by stable source IDs.

**Cross-encoder reranking.** We will rerank with the pretrained cross-encoder bge-reranker-v2-m3 [3], as a fixed off-the-shelf reranking component. It reads the question and each pooled paragraph together and outputs a relevance score, as in monoBERT reranking [16].

**Budgeted evidence selection.** A greedy selector will repeatedly add the paragraph with the highest gain among those that still fit in the budget  $B$ , until no paragraph fits. For a candidate paragraph  $p$  and the current selected set  $S$ , the gain is

$$\begin{aligned} \Delta(p \mid S) &= R(p) + \mu n_S(\text{sec}(p)). \\ R(p) &= \alpha r_{\text{ce}}(p) + \beta r_{\text{bm25}}(p) + \gamma r_{\text{dense}}(p), \end{aligned} \quad (1)$$

We will admit a candidate only if  $c(S \cup \{p\}) \leq B$ . The  $r$  terms are min-max normalized relevance scores over a fixed candidate pool, with  $\alpha, \beta, \gamma \geq 0$  and  $\alpha + \beta + \gamma = 1$ . We will count the serialized evidence text, source headers, and separators with the generator tokenizer and verify that the final context fits  $B$ . Fixed instructions and the question will be counted separately in total prompt tokens. The count  $n_S(\text{sec}(p))$  is the number of already selected paragraphs from  $p$ ’s supplied section (identified by its index, not just its title), and  $\mu \geq 0$  is the section-affinity weight. Token cost is a budget constraint, not part of the score. Setting  $\mu = 0$  recovers ranking by  $R(p)$  with budgeted packing: scan in score order, skip paragraphs that do not fit, and keep whole paragraphs. It recovers pure cross-encoder packing when  $\alpha = 1$  and  $\beta = \gamma = 0$ . The section-affinity selector is a heuristic with no optimality guarantee. Exact paragraph IDs will be kept for provenance. Similarity will be embedding cosine, with an empty-set redundancy penalty of zero. All methods will present selected paragraphs to the generator in source order.

**Why section affinity.** It is structure-aware expansion at paragraph level, like merging chunks into parent sections [15] or retrieving larger units [9, 22], but applied at selection time under a fixed budget, so it can be scored against gold paragraphs. We will tune  $\mu$  on dev with the cross-encoder in place and compare single- and multi-evidence questions. A zero affinity weight remains a valid choice if concentration does not help.

**Ablation rows.** We will keep two redundancy-aware selectors as ablation rows: relevance per token with max-similarity redundancy,  $(R(p) - \lambda \max_{s \in S} \text{sim}(p, s))/c(p)$ , which adapts the packing scores of Matrix Context [1] and SF-RAG [27]; and maximal marginal relevance,  $R(p) - \lambda \max_{s \in S} \text{sim}(p, s)$  [5]. These controls will test whether length preference or redundancy penalties improve packed evidence recall. We will not stop adaptively: adding a feasible paragraph without displacing selected evidence cannot decrease evidence recall, although it may affect precision or answer quality, and min-max normalized scores offer no absolute threshold.

**Generation with provenance.** The selected paragraphs will be passed, with their IDs, to one fixed LLM (Qwen2.5-7B-Instruct [19]). The system will return the answer together with the paper ID, section title, and exact text of every supplied paragraph. We will

request paragraph-ID citations and validate that cited IDs belong to the supplied context; this will supply provenance without assuming that citations prove correctness.

## 2.2 Original Implementation

We will implement BM25 indexing, exact dense search, RRF and convex fusion, candidate-pool control, section-aware selection, and the evaluation harness. We will reuse pretrained model inference and the official QASPER evaluator [7]; the packing trace is inspired by Matrix Context [1]. The original harness will implement mapping, budget checks, oracle ceilings, loss diagnostics, and paired statistics.

## 2.3 Why We Expect Gains over the Baselines

BM25 matches lexical identifiers, dense retrieval matches paraphrases, and fusion may recover complementary evidence [4]. A cross-encoder can improve relevance ranking [16]; section affinity will test whether evidence concentration adds value. Every budget table will include a gold-only oracle row. For each eligible annotation, it will find the largest feasible subset of mapped gold paragraphs under the same serialized token costs (by subset enumeration), then take the maximum over annotations. Full-evidence feasibility will be reported separately. Arbitrary gold-first ordering is not an optimal recall ceiling.

## 2.4 Related Work

The closest study is Singh et al. [24], which uses our two benchmarks and asks whether retrieval add-ons help once a strong reranker is present. We will extend it along the axes it leaves open: pairwise reranking, weaker first stages, the reranking pool size, gold evidence in the packed context, and the selection stage. Budgeted packing has been studied through context-budget diagnostics on QASPER [11] and coverage-based packing on multi-hop data [2]. We will add a stage-by-stage attribution and a structure-aware selector.

The overall pipeline shape (hybrid retrieval, reranking, budgeted greedy packing) already runs in Matrix Context [1]. It uses one packer for all methods, on a synthetic benchmark at one budget; we will borrow its trace of kept and dropped items for loss attribution. SF-RAG [27] combines structure, reranking, and budgeted assembly on QASPER, but uses an LLM for section selection and also reports evidence F1 and structural diagnostics. Our distinction is the controlled first-stage/pool study, budget sweep, and lightweight section-affinity selector, rather than being the first to evaluate evidence-aware assembly. Agentic systems such as VecTree-RAG [28] put an LLM inside retrieval; we will not.

## 3 IMPLEMENTATION PLAN

We will write the pipeline in Python and reuse only pretrained models and generic numeric libraries; every retrieval, fusion, selection, and evaluation step will be our own code. No RAG framework (LlamaIndex, LangChain) will be used. Table 1 separates what we will reuse from what we will build.

**LLM access and use.** We plan to serve Qwen2.5-7B-Instruct [19] locally with vLLM [12] on a compatible PACE ICE GPU, subject to allocation and a milestone feasibility check; one generator, one

**Table 1: Components we will reuse versus build.**

| Component                                                      | Reuse / build                               | Tools                                                                    |
|----------------------------------------------------------------|---------------------------------------------|--------------------------------------------------------------------------|
| Data loading, paragraph chunking, gold mapping                 | Build                                       | Python; public QASPER and MuSiQue releases                               |
| BM25 index and scoring                                         | Build                                       | Python inverted index, NumPy                                             |
| Dense embeddings                                               | Reuse model; build index and search         | sentence-transformers [20], all-MiniLM-L6-v2 [23], exact cosine in NumPy |
| Hybrid fusion (RRF, convex)                                    | Build                                       | Python                                                                   |
| Cross-encoder reranking                                        | Reuse model; build batching and pool logic  | bge-reranker-v2-m3 [3] via Hugging Face Transformers                     |
| Budgeted evidence selector                                     | Build                                       | Python, NumPy                                                            |
| Generation with provenance                                     | Reuse LLM; build prompt and citation output | vLLM [12], Qwen2.5-7B-Instruct [19]                                      |
| Evaluation harness, per-stage loss trace, bootstrap statistics | Build; reuse F1 evaluator                   | Python, pandas; official QASPER evaluator [7] for answer and evidence F1 |
| Latency and cost logging                                       | Build                                       | Python timers, scored pairs, prompt tokens, GPU runtime                  |

prompt, and greedy decoding will be shared by every configuration. No LLM will be used inside retrieval, which keeps retrieval cheap, deterministic, and comparable to the baselines.

## 4 EVALUATION PLAN

Our primary metric is gold-evidence recall of the packed context at a fixed token budget, i.e. of what the generator actually receives, with answer quality and system cost reported alongside.

### 4.1 Metrics of Success

Table 2 lists the metrics. For a question with a nonempty gold evidence set  $G$  and selected set  $S$  under budget  $B$ , evidence recall is  $|S \cap G|/|G|$  and all-evidence coverage is 1 if  $G \subseteq S$  and 0 otherwise, averaged over questions. QASPER questions often have several annotators with different evidence sets. We will score against the best-matching annotator, as VecTree-RAG does [28]. Because selecting the best annotator is optimistic, we will always report union-of-annotators recall next to it. For paragraph retrieval, an eligible question has at least one answerable, nonempty, text-only evidence annotation. We will report exclusions for unanswerable, empty-evidence, and figure/table-only cases, and score each metric against eligible annotations only. These will be text-evidence subset metrics; full-task answer F1 will include unanswerables. Subgroups will use the minimum deduplicated gold count across eligible annotations.

### 4.2 Benchmarks

**QASPER (primary).** We will follow QASPER’s per-paper setting [7]: the paper is known and we will retrieve among its paragraphs. An optional secondary pooled setting will retrieve across all papers of the split; this is where first-stage recall limits everything downstream. We will run it with and without the paper title prepended to each question, since the title alone can identify the paper, and report paper-level recall@ $k$  separately.

**MuSiQue-Ans (secondary).** MuSiQue [25] contains multi-hop questions whose context has 20 paragraphs, including annotated

**Table 2: Metrics we will report.**

| Category            | Metric                                                                          | What it answers                                                                |
|---------------------|---------------------------------------------------------------------------------|--------------------------------------------------------------------------------|
| Retrieval (primary) | Evidence recall at budget $B$                                                   | What fraction of gold paragraphs reached the packed context?                   |
| Retrieval (primary) | All-evidence coverage at $B$                                                    | Did every gold paragraph reach it?                                             |
| Retrieval           | Evidence-F1 [7]                                                                 | Comparable with prior QASPER work                                              |
| Retrieval           | Recall@ $k$ , MRR, nDCG@10                                                      | Ranking quality before selection; nDCG@10 is comparable with Singh et al. [24] |
| Retrieval           | Oracle recall at $B$                                                            | How much recall does the budget allow at all?                                  |
| Retrieval           | Loss attribution                                                                | Was the missing evidence absent from the pool or lost during selection?        |
| QA (secondary)      | Answer F1 (official QASPER evaluator); EM and F1 on MuSiQue                     | Does better evidence yield better answers?                                     |
| System              | Scored question-paragraph pairs, p50/p95 latency, prompt tokens and GPU runtime | What does each gain cost?                                                      |

supporting paragraphs for each hop; each question needs 2 to 4 supporting paragraphs. Its 20-paragraph context mirrors QASPER’s per-paper pool, and, as an optional extension, we will build a deduplicated pooled corpus from the evaluation contexts to test the first stage at scale. This will be our custom pooled setting, distinct from the supplied-context benchmark; we will report corpus size and preserve supporting-paragraph IDs.

Metrics will be supporting-paragraph recall and all-support coverage at budget  $B$ , plus answer EM and F1. Later-hop paragraphs are often reachable only through an earlier hop’s answer. Because our retrieval will be single-shot, we will report recall per hop. On MuSiQue we will run the RQ1 slice, which will make the comparison with Singh et al. [24] cover both of their homogeneous benchmarks.

### 4.3 Evaluation Sets and Statistics

Table 3 gives the splits. On QASPER we will evaluate every eligible test question; answer F1 will use a stratified subset and only BM25, dense, and the final proposed method at  $B = 1024$ . MuSiQue QA will use 200 fixed questions from its retrieval evaluation sample. Both QA samples will use a fixed seed and identical question IDs across methods.

Every method will see the same questions, so we will compare methods with a paired bootstrap [8] that resamples papers (questions about one paper are correlated) and report 95% confidence intervals. We will resample QASPER papers and MuSiQue questions with 10,000 paired bootstrap replicates. We will report effect sizes and 95% intervals. Planned comparisons against both baselines and the cross-encoder-only packer will use Holm correction within each benchmark; other grid comparisons will be descriptive. Interval widths and eligible sample counts will be measured rather than assumed. Full QASPER retrieval evaluation and 1,000 MuSiQue questions will support paired comparisons; subgroup results may be less precise. We will report subgroup counts (evidence count, answer type, hop count).

**Table 3: Tuning and evaluation sets.**

| Benchmark   | Tuning set                                 | Evaluation set                                                      | Selection                                                                        |
|-------------|--------------------------------------------|---------------------------------------------------------------------|----------------------------------------------------------------------------------|
| QASPER      | Dev split: 1,005 questions, 281 papers [7] | Test split: 1,451 questions, 416 papers [7]; all eligible questions | Full split for retrieval; 500 questions stratified by answer type for generation |
| MuSiQue-Ans | 300 questions sampled from train           | 1,000 of the 2,417 dev questions [25]                               | Stratified by hop count (2, 3, 4)                                                |

**Table 4: Ablation grid.**

| Axis               | Values                                                                                    |
|--------------------|-------------------------------------------------------------------------------------------|
| First stage        | BM25, dense, hybrid RRF, hybrid convex                                                    |
| Cross-encoder pool | 5, 10, 20, 30, whole paper (per-paper only); no reranker                                  |
| Selector           | top- $k$ packing; section affinity (Equation 1); per-token redundancy and MMR (ablations) |
| Selector features  | cross-encoder only; cross-encoder + BM25 + dense                                          |
| Token budget       | 256, 512, 1024, 2048 tokens                                                               |
| Setting            | per-paper (primary); pooled, with and without title                                       |

### 4.4 Ablation Grid

Table 4 lists the axes. We will run one slice per research question, not the full cross-product. **RQ1** will cross first stage with pool size at every budget, with top- $k$  packing. **RQ2** will cross selector with reranker on/off, budget, and evidence count, using the first stage selected on the tuning set from RQ1; the reranker-off cells will isolate the selector without cross-encoder scores, using first-stage relevance with fixed dev-tuned settings. **RQ3** will compare the two feature sets with the RQ2 selector fixed. We will cache scores for each question-paragraph pair to replay quality ablations. Query-time latency runs will bypass score caches, recompute query encoding and reranking with fixed batching, and count scored pairs for the actual pool  $N$ . Offline embedding/indexing cost will be reported separately; cached replay time is not serving latency.

## 5 TIMELINE AND DIVISION OF WORK

**By the milestone report.** By W11, we will complete QASPER and MuSiQue loading, gold mapping, both retrieval baselines, hybrid fusion, and the evaluation harness with oracle and bootstrap. We will report dev-set evidence diagnostics and baseline results, and demonstrate vLLM serving with provenance on a small sample.

**By the final report (W16).** Cross-encoder reranking with pool control; the section-affinity selector and its ablations; the full grid on the QASPER test split in the per-paper setting; optional pooled experiments after the core evaluations; MuSiQue evaluation; answer F1 on the generation subsets with latency and cost; loss attribution, error analysis, and the final report.

**Division of work.** Everyone will run ablations and write their part of the report.

- **Navneet:** data loading, chunking, gold mapping, and the evaluation harness and statistics; will also lead error analysis.
- **Jigyasa Hasija:** BM25, the dense index, hybrid fusion, and the cross-encoder pipeline; will also run the pooled-setting experiments.

- **Joseph Souaiby:** the budgeted selector, generation with provenance, and LLM serving; will also handle latency and cost logging.

## REFERENCES

- [1] agent-matrix. 2026. Matrix Context: Mixture-of-Contexts Retrieval (MoC-RAG). GitHub repository. <https://github.com/agent-matrix/matrix-context>
- [2] Ananto Nayan Bala. 2026. What Survives Into Context: A Diagnostic for Budget-Constrained Multi-Hop RAG and When Submodular Evidence Packing Improves It. arXiv preprint arXiv:2607.00725. <https://arxiv.org/abs/2607.00725>
- [3] Beijing Academy of Artificial Intelligence. [n.d.]. BAAI/bge-reranker-v2-m3: Model Card. Hugging Face. <https://huggingface.co/BAAI/bge-reranker-v2-m3> Accessed September 30, 2026.
- [4] Sebastian Bruch, Siyu Gai, and Amir Ingber. 2023. An Analysis of Fusion Functions for Hybrid Retrieval. *ACM Transactions on Information Systems* 42, 1 (2023). <https://doi.org/10.1145/3596512>
- [5] Jaime Carbonell and Jade Goldstein. 1998. The Use of MMR, Diversity-Based Reranking for Reordering Documents and Producing Summaries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*. 335–336.
- [6] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Buettcher. 2009. Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*. 758–759.
- [7] Pradeep Dasigi, Kyle Lo, Iz Beltagy, Arman Cohan, Noah A. Smith, and Matt Gardner. 2021. A Dataset of Information-Seeking Questions and Answers Anchored in Research Papers. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*. 4599–4610. <https://aclanthology.org/2021.naacl-main.365/>
- [8] Bradley Efron and Robert J. Tibshirani. 1993. *An Introduction to the Bootstrap*. Chapman & Hall.
- [9] Ziyang Jiang, Xueguang Ma, and Wenhui Chen. 2024. LongRAG: Enhancing Retrieval-Augmented Generation with Long-context LLMs. arXiv preprint arXiv:2406.15319. <https://arxiv.org/abs/2406.15319>
- [10] Vladimir Karpukhin, Barlas Öguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 6769–6781.
- [11] Seowan Kim and Jinhong Yang. 2026. Evidence Retention Is Not Sufficient for Answerability: A Diagnostic Study of Context Budgeting for Small Language Model RAG. In *Proceedings of the 14th International Conference on Advances in Information Technology (IAIT '26)*. ACM. <https://doi.org/10.1145/3816713.3819786>
- [12] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP)*.
- [13] Nathanaël Langlois. 2026. SCAR: Semantic Continuity-Aware Retrieval for Efficient Context Expansion in RAG. arXiv preprint arXiv:2606.16661. <https://arxiv.org/abs/2606.16661>
- [14] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 33.
- [15] LlamaIndex. 2026. Auto Merging Retriever. LlamaIndex documentation, accessed September 2026. [https://developers.llamaindex.ai/python/framework/integrations/retrievers/auto\\_merging\\_retriever/](https://developers.llamaindex.ai/python/framework/integrations/retrievers/auto_merging_retriever/)
- [16] Rodrigo Nogueira and Kyunghyun Cho. 2019. Passage Re-ranking with BERT. arXiv preprint arXiv:1901.04085. <https://arxiv.org/abs/1901.04085>
- [17] Chao Peng, Bin Wang, Zhilei Long, and Jinfang Sheng. 2025. AdaGRS: Adaptive Greedy Context Selection via Redundancy-Aware Scoring for Token-Budgeted RAG. arXiv preprint arXiv:2512.25052. <https://arxiv.org/abs/2512.25052>
- [18] Jayr Pereira, Robson Fidalgo, Roberto Lotufo, and Rodrigo Nogueira. 2023. Visconde: Multi-document QA with GPT-3 and Neural Reranking. In *Advances in Information Retrieval (ECIR)*. <https://arxiv.org/abs/2212.09656>
- [19] Qwen Team, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, et al. 2024. Qwen2.5 Technical Report. arXiv preprint arXiv:2412.15115. <https://arxiv.org/abs/2412.15115>
- [20] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP)*. 3982–3992.
- [21] Stephen Robertson and Hugo Zaragoza. 2009. The Probabilistic Relevance Framework: BM25 and Beyond. *Foundations and Trends in Information Retrieval* 3, 4 (2009), 333–389.
- [22] Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. 2024. RAPTOR: Recursive Abstractive Processing for Tree-Organized Retrieval. In *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/2401.18059>
- [23] Sentence Transformers. [n.d.]. sentence-transformers/all-MiniLM-L6-v2: Model Card. Hugging Face. <https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2> Accessed September 30, 2026.
- [24] Sadanand Singh, Allam Reddy, and Manan Chopra. 2026. Beyond the Reranker: Do RAG Retrieval Enhancements Help Once a Strong Reranker Is Present? arXiv preprint arXiv:2606.28367. <https://arxiv.org/abs/2606.28367>
- [25] Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. 2022. MuSiQue: Multihop Questions via Single-hop Question Composition. *Transactions of the Association for Computational Linguistics* 10 (2022), 539–554. <https://aclanthology.org/2022.tacl-1.31/>
- [26] Shuai Wang, Shengyao Zhuang, and Guido Zuccon. 2021. BERT-based Dense Retrievers Require Interpolation with BM25 for Effective Passage Retrieval. In *Proceedings of the 2021 ACM SIGIR International Conference on Theory of Information Retrieval (ICTIR)*. <https://doi.org/10.1145/3471158.3472233>
- [27] Rui Yu, Tianyi Wang, Ruixia Liu, and Yinglong Wang. 2026. SF-RAG: Structure-Fidelity Retrieval-Augmented Generation for Academic Question Answering. arXiv preprint arXiv:2602.13647. <https://arxiv.org/abs/2602.13647>
- [28] Xinyan Zhong, Yuwei Shi, Yuqi Wei, Chen Shen, Tianhang Zhou, and Zhenghao Wu. 2026. VecTree-RAG: An Agentic Retrieval-Augmented Generation Framework Combining Vector and Tree Retrieval for Efficiency and Accuracy. arXiv preprint arXiv:2607.23006. <https://arxiv.org/abs/2607.23006>